

Rolling window features: part 3

Window features

Rolling window features: Pandas

pandas.DataFrame.rolling

DataFrame.rolling(window, min_periods=None, center=False, win_type=None, on=None, axis=0, closed=None)

[\[source\]](#)

Provide rolling window calculations.

Parameters: **window** : *int, offset, or BaseIndexer subclass*

Size of the moving window. This is the number of observations used for calculating the statistic. Each window will be a fixed size.

If its an offset then this will be the time period of each window. Each window will be a variable sized based on the observations included in the time-period. This is only valid for datetimelike indexes.

If a BaseIndexer subclass is passed, calculates the window boundaries based on the defined `get_window_bounds` method. Additional rolling keyword arguments, namely *min_periods*, *center*, and *closed* will be passed to `get_window_bounds`.

min_periods : *int, default None*

Minimum number of observations in window required to have a value (otherwise result is NA). For a window that is specified by an offset, *min_periods* will default to 1. Otherwise, *min_periods* will default to the size of the window.

center : *bool, default False*

Set the labels at the center of the window.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3).agg(["mean", "min"])
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	NaN	NaN
1992-03-01	150930.333333	146376.0
1992-04-01	156694.666667	147079.0
1992-05-01	164357.666667	159336.0

The `rolling` method by default assigns the rolling statistics to the edge of the window.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3).agg(["mean", "min"])
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	NaN	NaN
1992-03-01	150930.333333	146376.0
1992-04-01	156694.666667	147079.0
1992-05-01	164357.666667	159336.0

For forecasting we want features which only use information that we will have at predict time (i.e, the past). So we want to shift the output of the row down by one to avoid data leakage.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3).agg(["mean", "min"]).shift(periods=1)
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	NaN	NaN
1992-03-01	NaN	NaN
1992-04-01	150930.333333	146376.0
1992-05-01	156694.666667	147079.0

For forecasting we want features which only use information that we will have at predict time (i.e, the past). So we want to shift the output of the row down by one to avoid data leakage.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3, min_periods=1).agg(["mean", "min"]).shift(periods=1)
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	146376.000000	146376.0
1992-03-01	146727.500000	146376.0
1992-04-01	150930.333333	146376.0
1992-05-01	156694.666667	147079.0

The ``min_periods`` argument allows us to use smaller windows at the edges.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3, min_periods=1).agg(["mean", "min"]).shift(periods=1)
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	146376.000000	146376.0
1992-03-01	146727.500000	146376.0
1992-04-01	150930.333333	146376.0
1992-05-01	156694.666667	147079.0

The ``min_periods`` argument allows us to use smaller windows at the edges.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3, min_periods=1).agg(["mean", "min"]).shift(periods=1)
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	146376.000000	146376.0
1992-03-01	146727.500000	146376.0
1992-04-01	150930.333333	146376.0
1992-05-01	156694.666667	147079.0

The ``min_periods`` argument allows us to use smaller windows at the edges.

Rolling window features: Pandas

```
df.head()
```

	y
ds	
1992-01-01	146376
1992-02-01	147079
1992-03-01	159336
1992-04-01	163669
1992-05-01	170068

```
df["y"].rolling(window=3, min_periods=1).agg(["mean", "min"]).shift(periods=1)
```

	mean	min
ds		
1992-01-01	NaN	NaN
1992-02-01	146376.000000	146376.0
1992-03-01	146727.500000	146376.0
1992-04-01	150930.333333	146376.0
1992-05-01	156694.666667	147079.0

The ``min_periods`` argument allows us to use smaller windows at the edges.

Rolling window features: Feature-engine

```
from feature_engine.timeseries.forecasting import WindowFeatures

# Create transformer for window features
transformer = WindowFeatures(variables=['y'],
                             functions=['mean', 'std'], # Stats
                             window=[1, 3, 6], # Window sizes
                             freq='1MS')
transformer.fit_transform(df)
```

	y	y_window_1_mean	y_window_1_std	y_window_3_mean	y_window_3_std	y_window_6_mean	y_window_6_std
ds							
2016-01-01	400928	518253.00	NaN	469239.67	42447.39	458781.00	30709.03
2016-02-01	413554	400928.00	NaN	454562.67	59305.36	449317.33	38790.92
2016-03-01	460093	413554.00	NaN	444245.00	64402.97	442186.33	41105.40
2016-04-01	450935	460093.00	NaN	424858.33	31160.32	447049.00	41231.17
2016-05-01	471421	450935.00	NaN	441527.33	24654.57	448045.00	41242.76

Rolling window features: Feature-engine

```
from feature_engine.timeseries.forecasting import WindowFeatures

# Create transformer for window features
transformer = WindowFeatures(variables=['y'],
                             functions=['mean', 'std'], # Stats
                             window=[1, 3, 6], # Window sizes
                             freq='1MS')
transformer.fit_transform(df)
```

	y	y_window_1_mean	y_window_1_std	y_window_3_mean	y_window_3_std	y_window_6_mean	y_window_6_std
ds							
2016-01-01	400928	518253.00	NaN	469239.67	42447.39	458781.00	30709.03
2016-02-01	413554	400928.00	NaN	454562.67	59305.36	449317.33	38790.92
2016-03-01	460093	413554.00	NaN	444245.00	64402.97	442186.33	41105.40
2016-04-01	450935	460093.00	NaN	424858.33	31160.32	447049.00	41231.17
2016-05-01	471421	450935.00	NaN	441527.33	24654.57	448045.00	41242.76

Rolling window features: Feature-engine

```
from feature_engine.timeseries.forecasting import WindowFeatures

# Create transformer for window features
transformer = WindowFeatures(variables=['y'],
                             functions=['mean', 'std'], # Stats
                             window=[1, 3, 6], # Window sizes
                             freq='1MS')
transformer.fit_transform(df)
```

	y	y_window_1_mean	y_window_1_std	y_window_3_mean	y_window_3_std	y_window_6_mean	y_window_6_std
ds							
2016-01-01	400928	518253.00	NaN	469239.67	42447.39	458781.00	30709.03
2016-02-01	413554	400928.00	NaN	454562.67	59305.36	449317.33	38790.92
2016-03-01	460093	413554.00	NaN	444245.00	64402.97	442186.33	41105.40
2016-04-01	450935	460093.00	NaN	424858.33	31160.32	447049.00	41231.17
2016-05-01	471421	450935.00	NaN	441527.33	24654.57	448045.00	41242.76

Rolling window features: sktime

```
from sktime.transformations.series.summarize import WindowSummarizer

transformer = WindowSummarizer(
    lag_feature={
        "lag": [1, 2] # Create lag features
        "mean": [[1, 3], [3, 6]], # [[lag, window size], ...]
        "std": [[1, 3]],
    },
    target_cols=["y"],
)

transformer.fit_transform(df)
```

Rolling window features: sktime

		<u>y</u>					
ds			<u>y_lag_1</u>	y_lag_2	y_mean_1_3	y_mean_3_6	y_std_1_4
1992-01-01	146376		NaN	NaN	NaN	NaN	NaN
1992-02-01	147079		146376.00	NaN	NaN	NaN	NaN
1992-03-01	159336		147079.00	146376.00	NaN	NaN	NaN
1992-04-01	163669		159336.00	147079.00	150930.33	NaN	NaN
1992-05-01	170068		163669.00	159336.00	156694.67	NaN	8716.56
...
2016-01-01	400928		518253.00	444507.00	469239.67	450128.33	39602.33
2016-02-01	413554		400928.00	518253.00	454562.67	447110.33	48660.13
2016-03-01	460093		413554.00	400928.00	444245.00	458781.00	52584.96
2016-04-01	450935		460093.00	413554.00	424858.33	449317.33	53178.48
2016-05-01	471421		450935.00	460093.00	441527.33	442186.33	28588.60

Rolling window features: sktime

		y	y_lag_1	<u>y_lag_2</u>	y_mean_1_3	y_mean_3_6	y_std_1_4
ds		ds					
1992-01-01	146376	1992-01-01	NaN	NaN	NaN	NaN	NaN
1992-02-01	147079	1992-02-01	146376.00	NaN	NaN	NaN	NaN
1992-03-01	159336	1992-03-01	147079.00	146376.00	NaN	NaN	NaN
1992-04-01	163669	1992-04-01	159336.00	147079.00	150930.33	NaN	NaN
1992-05-01	170068	1992-05-01	163669.00	159336.00	156694.67	NaN	8716.56
...
2016-01-01	400928	2016-01-01	518253.00	444507.00	469239.67	450128.33	39602.33
2016-02-01	413554	2016-02-01	400928.00	518253.00	454562.67	447110.33	48660.13
2016-03-01	460093	2016-03-01	413554.00	400928.00	444245.00	458781.00	52584.96
2016-04-01	450935	2016-04-01	460093.00	413554.00	424858.33	449317.33	53178.48
2016-05-01	471421	2016-05-01	450935.00	460093.00	441527.33	442186.33	28588.60

Rolling window features: sktime

y		y_lag_1	y_lag_2	<u>y_mean_1_3</u>	y_mean_3_6	y_std_1_4
ds		ds				
1992-01-01	146376	1992-01-01	NaN	NaN	NaN	NaN
1992-02-01	147079	1992-02-01	146376.00	NaN	NaN	NaN
1992-03-01	159336	1992-03-01	147079.00	146376.00	NaN	NaN
1992-04-01	163669	1992-04-01	159336.00	147079.00	150930.33	NaN
1992-05-01	170068	1992-05-01	163669.00	159336.00	156694.67	NaN
...
2016-01-01	400928	2016-01-01	518253.00	444507.00	469239.67	450128.33
2016-02-01	413554	2016-02-01	400928.00	518253.00	454562.67	447110.33
2016-03-01	460093	2016-03-01	413554.00	400928.00	444245.00	458781.00
2016-04-01	450935	2016-04-01	460093.00	413554.00	424858.33	449317.33
2016-05-01	471421	2016-05-01	450935.00	460093.00	441527.33	442186.33

Rolling window features: sktime

		y						
ds			y_lag_1	y_lag_2	y_mean_1_3	<u>y_mean_3_6</u>	y_std_1_4	
1992-01-01	146376		NaN	NaN	NaN	NaN	NaN	
1992-02-01	147079		146376.00	NaN	NaN	NaN	NaN	
1992-03-01	159336		147079.00	146376.00	NaN	NaN	NaN	
1992-04-01	163669		159336.00	147079.00	150930.33	NaN	NaN	
1992-05-01	170068		163669.00	159336.00	156694.67	NaN	8716.56	
...	
2016-01-01	400928		518253.00	444507.00	469239.67	450128.33	39602.33	
2016-02-01	413554		400928.00	518253.00	454562.67	447110.33	48660.13	
2016-03-01	460093		413554.00	400928.00	444245.00	458781.00	52584.96	
2016-04-01	450935		460093.00	413554.00	424858.33	449317.33	53178.48	
2016-05-01	471421		450935.00	460093.00	441527.33	442186.33	28588.60	

Summary

New features can be created by applying a rolling window to the target or other features.

The window features still need to be lagged to ensure there is no data leakage.

Multiple different window sizes could be helpful. The seasonal period and different time scales can be a good starting point.

Mean and standard deviation are common. Use feature selection methods for more statistics.